



COMP - 6521
Advanced Database Technology and Applications

Project Report
On

Lab Assignment 1: TPMMS & Duplicate Elimination

Professor
Dr. Nematollaah Shiri

Team Members

Student Name	Student Id	Email Id
Yash Pandya	40119272	y_pandy@encs.concordia.ca
Himen Sidhpura	40091993	h_sidhpu@encs.concordia.ca
Sucheta Sudhakumari	40080543	s_ijaya@encs.concordia.ca

Index

1. Steps to run the program	3
2. Program Description	3
3. Experiment Results	3
4. Architecture Diagrams	6
5. Algorithm	7
6. Technical Details	7
7. Coding Standard	7
8. Class Description	8
9. Group Member Contribution	8
10. Results	9

1. Steps to run the program:

- Place the data sets in the 'InputFiles' folder and set up the eclipse environment.
- Setup main memory size.
- Run the program with 'ProgramController.java' class.
- The program will read the input files based on memory size, sort them into sublists, merge that sublists into final sorted relation while eliminating the duplicates.
- The sub lists can be seen in the 'blocks' folder.
- The output which is the final sorted list without duplicates will be generated in the output files folder.

2. Program Description:

- In Phase 1, the program begins by clearing the output and sublist folders. Followed by this the program takes the data sets T1(sample1.txt) and T2(sample2.txt) one after the other and sorts them. The sub-lists created in phase 1 can be seen in the blocks folder.
- In phase 2, the sublists are read into the 2 input buffers iteratively and compared against each other to find the smallest one which will then be written into the output buffer. Elimination of duplicate tuples takes place amidst this comparison. Once there is no more sublists remaining the final output is written back to the disk.

3. Experiment Results:

The experiment was conducted by changing the allocated main memory size the result of which has been consolidated below:

Experiment 1:

Main Memory Size = 5 MB

Tuple Count	20,00,064
Number of Disk IO for phase 1	1,00,004
Time to sort T1	3.633 sec
Time to sort T2	4.008 sec
Number of Blocks created	50,001
Number of Disk IO for phase 2	1,46,581
Time to merge Sublists	5.015 sec
Total program execution time	12.656 sec

Main Memory Size = 10 MB

Tuple Count	20,00,064
Number of Disk IO for phase 1	1,00,004
Time to sort T1	3.103 sec
Time to sort T2	3.626 sec
Number of Blocks created	50,001
Number of Disk IO for phase 2	1,06,230
Time to merge Sublists	2.831 sec
Total program execution time	9.56 sec

Main Memory Size = 20 MB

Tuple Count	20,00,064
Number of Disk IO for phase 1	1,00,004
Time to sort T1	2.494 sec
Time to sort T2	2.404 sec
Number of Blocks created	50,001
Number of Disk IO for phase 2	77,350
Time to merge Sublists	1.745 sec
Total program execution time	6.642 sec

Experiment 2:**Main Memory Size = 5 MB**

Tuple Count	12,00,064
Number of Disk IO for phase 1	60,004
Time to sort T1	2.506 sec
Time to sort T2	2.239 sec
Number of Blocks created	30,001
Number of Disk IO for phase 2	88,210
Time to merge Sublists	2.918 sec

Total program execution time	7.663 sec
------------------------------	-----------

Main Memory Size = 10 MB

Tuple Count	12,00,064
Number of Disk IO for phase 1	60,004
Time to sort T1	2.011 sec
Time to sort T2	1.933 sec
Number of Blocks created	30,001
Number of Disk IO for phase 2	63,738
Time to merge Sublists	1.839 sec
Total program execution time	5.783 sec

Main Memory Size = 20 MB

Tuple Count	12,00,064
Number of Disk IO for phase 1	60,004
Time to sort T1	1.55 sec
Time to sort T2	1.5 sec
Number of Blocks created	30,001
Number of Disk IO for phase 2	46,384
Time to merge Sublists	0.963 sec
Total program execution time	4.013 sec

4. Architecture Diagram:

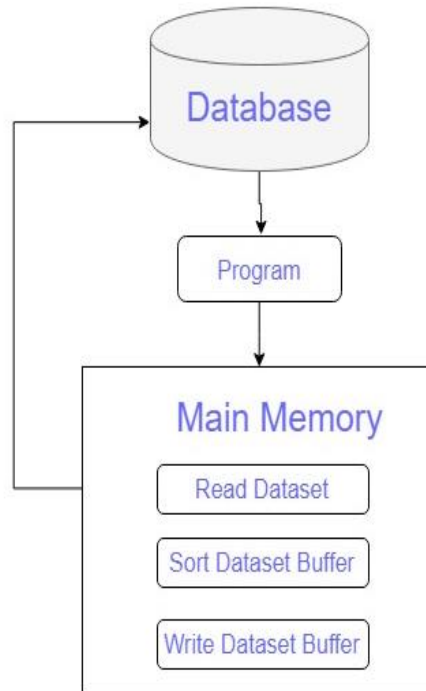


Figure 1: Phase 1

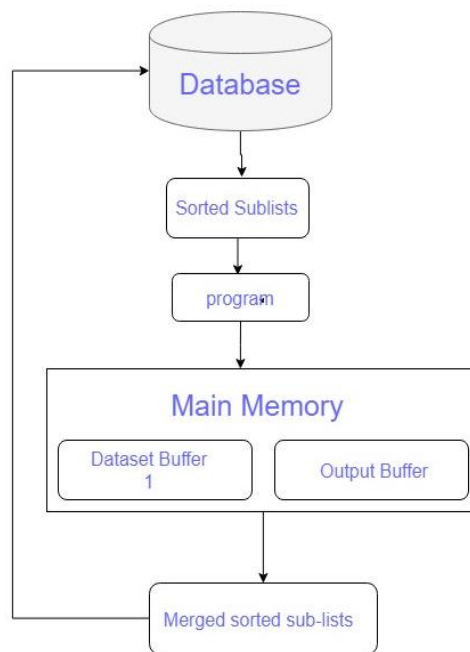


Figure 2: Phase 2

5. Algorithms:

PHASE - 1 (TPMMS Sort):

- Start.
- Allocate main memory to read sublists.
- Calculate the number of tuples that can be read into the input buffer in 1 iteration.
- Reset output folder and sublist folder.
- Load the input buffer.
- Update the count of I/O operation.
- Apply Quick sort to sort the tuples within input buffer.
- Write the sorted list to the sublist folder.
- Update the count of I/O operations.
- Update the count of sublists created.
- Repeat the above steps from 5 to 10 for calculated number of memory fills.
- Calculate the time taken to sort the relation.
- End.

PHASE - 2 (TPMMS Sort):

- Start
- Create array of input buffers for sublists and divide that 70% of available main memory space among sublist buffers equally.
- Create output buffer with 30% of available main memory space.
- Load the input sublist buffers and get array length.
- Read the first tuple from all the sublist buffers and find the smallest one.
- Swap that buffer having the smallest tuple with the 0th position buffer.
- Pop out smallest tuple from 0th position buffer and put into output buffer, and check refill is required for that buffer. If yes, refill the buffer.

6. Technical Detail:

The program executes in two phases

- Phase-1 involves reading input file data and saving sorted data in sublist.
- Phase-2 includes the collection of sublist records and the merging of the sorted list into the output file by removing the duplicate tuple and contains tuple with the latest date.

The above process is followed for memory size of 10 MB and 20 MB, then the execution period is measured for executing the entire operation that involves the sorting and the duplicate removal in the merging.

7. Coding Standards:

The most general coding conventions were followed while the codes were developed as follows,

- The class name begins with an uppercase word.
- E.g.: ProgramController.java
- Constants are called with characters in the upper case
- The variable name is descriptive and is rendered in lower case including a capital letter to separate words.
- The procedure name begins with a lowercase character and uses the uppercase characters to separate words.

8. Class Description:

Constants.java: This class stores the constant values required for program execution like file IO paths, block size etc.

PhaseOne.java: This class has the methods for reading tuples into main memory, sorting them and creating the sublists. It also handles the IO calculation for phase one.

PhaseTwo.java: This class has the methods for combining the sublists into a final sorted list, while eliminating duplicates.

ProgramController.java: This class has the main method from where program execution starts. It builds and clears the output directories and prints the timing outputs. It also creates the objects for other classes and trigger their execution.

QuickSort.java: This class performs the quicksort algorithm on the blocks read into memory.

9. Group Member Contribution:

10. Results:

```
*****Cleaning Directory*****
Block Directory Deleted :- true
Block Directory Created :- true
Output Directory Deleted :- true
Output Directory Created :- true
Diretory Cleaned
*****TPMMS Console*****
Memory Size : 5
Tupple Size : 100
*****Phase 1 for T1*****
Time taken by Phase 1 for T1 : 2506ms (2.506sec)
Records in T1 : 600032
Block for T1 : 15001
*****Phase 1 for T2*****
Time taken by Phase 1 for T2 : 2239ms (2.239sec)
Records in T2 : 600032
Block for T2 : 15001
*****Phase 1 Overview*****
Total number of records 1200064
Total number of Block 30001
Sorted Disk IO 60004
*****Phase 2*****
Phase 2 merging time iteration 0 : 1588ms(~approx 1.588sec)
Phase 2 merging time iteration 1 : 664ms(~approx 0.664sec)
Phase 2 merging time iteration 2 : 333ms(~approx 0.333sec)
Phase 2 merging time iteration 3 : 163ms(~approx 0.163sec)
Phase 2 merging time iteration 4 : 87ms(~approx 0.087sec)
Phase 2 merging time iteration 5 : 48ms(~approx 0.048sec)
Phase 2 merging time iteration 6 : 24ms(~approx 0.024sec)
Phase 2 merging time iteration 7 : 12ms(~approx 0.012sec)
Phase 2 Time : 2918ms (2.918 sec)
Merge Phase IO of I/O :88210
*****Output Overview*****
Total time Phase 1 & Phase 2 : 7663ms
Total time Phase 1 & Phase 2 : 7.663 sec
Total Number of I/O : 148214
```

Figure 3: 5 MB – 12,00,064 Tuple Count

```

<terminated> ProgramController [Java Application] C:\Program Files\Java\jre\bin\javaw.exe (Feb 16, 2020, 4:53:58 PM)
*****Cleaning Directory*****
Block Directory Deleted :- true
Block Directory Created :- true
Output Directory Deleted :- true
Output Directory Created :- true
Diretory Cleaned
*****TPMMS Console*****
Memory Size : 5
Tuple Size : 100
*****Phase 1 for T1*****
Time taken by Phase 1 for T1 : 3633ms (3.633sec)
Records in T1 : 1000032
Block for T1 : 25001
*****Phase 1 for T2*****
Time taken by Phase 1 for T2 : 4008ms (4.008sec)
Records in T2 : 1000032
Block for T2 : 25001
*****Phase 1 Overview*****
Total number of records 2000064
Total number of Block 50001
Sorted Disk IO 100004
*****Phase 2*****
Phase 2 merging time iteration 0 : 2550ms(~approx 2.55sec)
Phase 2 merging time iteration 1 : 1389ms(~approx 1.389sec)
Phase 2 merging time iteration 2 : 539ms(~approx 0.539sec)
Phase 2 merging time iteration 3 : 258ms(~approx 0.258sec)
Phase 2 merging time iteration 4 : 130ms(~approx 0.13sec)
Phase 2 merging time iteration 5 : 73ms(~approx 0.073sec)
Phase 2 merging time iteration 6 : 41ms(~approx 0.041sec)
Phase 2 merging time iteration 7 : 23ms(~approx 0.023sec)
Phase 2 merging time iteration 8 : 12ms(~approx 0.012sec)
Phase 2 Time : 5015ms (5.015 sec)
Merge Phase IO of I/O :146581
*****Output Overview*****
Total time Phase 1 & Phase 2 : 12656ms
Total time Phase 1 & Phase 2 : 12.656 sec
Total Number of I/O : 246585

```

Figure 4: 5MB – 20,00,064 Tuple Counts

```

*****Cleaning Directory*****
Block Directory Deleted :- true
Block Directory Created :- true
Output Directory Deleted :- true
Output Directory Created :- true
Diretory Cleaned
*****TPMS Console*****
Memory Size : 9
Tuple Size : 100
*****Phase 1 for T1*****
Time taken by Phase 1 for T1 : 2011ms (2.011sec)
Records in T1 : 600032
Block for T1 : 15001
*****Phase 1 for T2*****
Time taken by Phase 1 for T2 : 1933ms (1.933sec)
Records in T2 : 600032
Block for T2 : 15001
*****Phase 1 Overview*****
Total number of records 1200064
Total number of Block 30001
Sorted Disk IO 60004
*****Phase 2*****
Phase 2 merging time iteration 0 : 1002ms(~approx 1.002sec)
Phase 2 merging time iteration 1 : 342ms(~approx 0.342sec)
Phase 2 merging time iteration 2 : 177ms(~approx 0.177sec)
Phase 2 merging time iteration 3 : 240ms(~approx 0.24sec)
Phase 2 merging time iteration 4 : 45ms(~approx 0.045sec)
Phase 2 merging time iteration 5 : 23ms(~approx 0.023sec)
Phase 2 merging time iteration 6 : 10ms(~approx 0.01sec)
Phase 2 Time : 1839ms (1.839 sec)
Merge Phase IO of I/O :63738
*****Output Overview*****
Total time Phase 1 & Phase 2 : 5783ms
Total time Phase 1 & Phase 2 : 5.783 sec
Total Number of I/O : 123742

```

Figure 5: 10MB – 12,00,064 Tuple Count


```

total number of I/O : 1 (phaseOneIO + phaseTwoIO) + phaseOneIO + phaseTwoIO
<terminated> ProgramController [Java Application] C:\Program Files\Java\jre\bin\javaw.exe (Feb 16, 2020, 4:51:40 PM)
*****Cleaning Directory*****
Block Directory Deleted :- true
Block Directory Created :- true
Output Directory Deleted :- true
Output Directory Created :- true
Directory Cleaned
*****TPMMS Console*****
Memory Size : 9
Tuple Size : 100
*****Phase 1 for T1*****
Time taken by Phase 1 for T1 : 3103ms (3.103sec)
Records in T1 : 1000032
Block for T1 : 25001
*****Phase 1 for T2*****
Time taken by Phase 1 for T2 : 3626ms (3.626sec)
Records in T2 : 1000032
Block for T2 : 25001
*****Phase 1 Overview*****
Total number of records 2000064
Total number of Block 50001
Sorted Disk IO 100004
*****Phase 2*****
Phase 2 merging time iteration 0 : 1522ms(~approx 1.522sec)
Phase 2 merging time iteration 1 : 742ms(~approx 0.742sec)
Phase 2 merging time iteration 2 : 277ms(~approx 0.277sec)
Phase 2 merging time iteration 3 : 143ms(~approx 0.143sec)
Phase 2 merging time iteration 4 : 75ms(~approx 0.075sec)
Phase 2 merging time iteration 5 : 40ms(~approx 0.04sec)
Phase 2 merging time iteration 6 : 21ms(~approx 0.021sec)
Phase 2 merging time iteration 7 : 11ms(~approx 0.011sec)
Phase 2 Time : 2831ms (2.831 sec)
Merge Phase IO of I/O : 106230
*****Output Overview*****
Total time Phase 1 & Phase 2 : 9560ms
Total time Phase 1 & Phase 2 : 9.56 sec
Total Number of I/O : 206234

```

Figure 6: 10MB – 20,00,064 Tuple Count

```

*****Cleaning Directory*****
Block Directory Deleted :- true
Block Directory Created :- true
Output Directory Deleted :- true
Output Directory Created :- true
Diretory Cleaned
*****TPMMS Console*****
Memory Size : 19
Tuple Size : 100
*****Phase 1 for T1*****
Time taken by Phase 1 for T1 : 1550ms (1.55sec)
Records in T1 : 600032
Block for T1 : 15001
*****Phase 1 for T2*****
Time taken by Phase 1 for T2 : 1500ms (1.5sec)
Records in T2 : 600032
Block for T2 : 15001
*****Phase 1 Overview*****
Total number of records 1200064
Total number of Block 30001
Sorted Disk IO 60004
*****Phase 2*****
Phase 2 merging time iteration 0 : 676ms(~approx 0.676sec)
Phase 2 merging time iteration 1 : 143ms(~approx 0.143sec)
Phase 2 merging time iteration 2 : 73ms(~approx 0.073sec)
Phase 2 merging time iteration 3 : 42ms(~approx 0.042sec)
Phase 2 merging time iteration 4 : 20ms(~approx 0.02sec)
Phase 2 merging time iteration 5 : 10ms(~approx 0.01sec)
Phase 2 Time : 963ms (0.963 sec)
Merge Phase IO of I/O :46384
*****Output Overview*****
Total time Phase 1 & Phase 2 : 4013ms
Total time Phase 1 & Phase 2 : 4.013 sec
Total Number of I/O : 106388

```

Figure 7: 20MB – 12,00,064 Tuple Count

```

<terminated> ProgramController [Java Application] C:\Program Files\Java\jre\bin\javaw.exe (Feb 16, 2020, 4:53:09 PM)
*****Cleaning Directory*****
Block Directory Deleted :- true
Block Directory Created :- true
Output Directory Deleted :- true
Output Directory Created :- true
Diretory Cleaned
*****TPMMS Console*****
Memory Size : 19
Tuple Size : 100
*****Phase 1 for T1*****
Time taken by Phase 1 for T1 : 2494ms (2.494sec)
Records in T1 : 1000032
Block for T1 : 25001
*****Phase 1 for T2*****
Time taken by Phase 1 for T2 : 2403ms (2.404sec)
Records in T2 : 1000032
Block for T2 : 25001
*****Phase 1 Overview*****
Total number of records 2000064
Total number of Block 50001
Sorted Disk IO 100004
*****Phase 2*****
Phase 2 merging time iteration 0 : 1256ms(~approx 1.256sec)
Phase 2 merging time iteration 1 : 236ms(~approx 0.237sec)
Phase 2 merging time iteration 2 : 128ms(~approx 0.128sec)
Phase 2 merging time iteration 3 : 63ms(~approx 0.063sec)
Phase 2 merging time iteration 4 : 35ms(~approx 0.035sec)
Phase 2 merging time iteration 5 : 18ms(~approx 0.018sec)
Phase 2 merging time iteration 6 : 9ms(~approx 0.009sec)
Phase 2 Time : 1745ms (1.745 sec)
Merge Phase IO of I/O : 77350
*****Output Overview*****
Total time Phase 1 & Phase 2 : 6642ms
Total time Phase 1 & Phase 2 : 6.642 sec
Total Number of I/O : 177354

```

Figure 8: 10MB – 20,00,064 Tuple Count

References

1. <https://www.geeksforgeeks.org/java-program-for-quicksort/>
2. <https://examples.javacodegeeks.com/core-java/nio/bytebuffer/write-append-to-file-with-byte-buffer/>
3. <https://www.geeksforgeeks.org/bytebuffer-get-method-in-java-with-examples/>
4. http://rosettacode.org/wiki/Binary_search#Java
5. https://en.wikipedia.org/wiki/Sorting_algorithm
6. <https://crunchify.com/increase-eclipse-memory-size-to-avoid-oom-on-startup/>

7. https://github.com/sagarveta/ADB_Project_1_TPMMS
8. <https://howtodoinjava.com/java/io/how-to-check-if-file-exists-in-java/>
9. <https://www.javadevjournal.com/java/java-copy-file-directory/>
10. <https://mkyong.com/java/how-to-delete-directory-in-java/>
11. <https://www.tutorialspoint.com/how-to-create-a-new-directory-by-using-file-object-in-java>
12. https://www.tutorialspoint.com/java/io/file_isfile.htm
13. <http://www.mathcs.emory.edu/~cheung/Courses/554/Syllabus/4-query-exec/2-pass=TPMMS.html>
14. <http://www.mathcs.emory.edu/~cheung/Courses/554/Syllabus/4-query-exec/TPMMS=join2.html>